

Summer Data Scientist Data Assessment

Crime and Education Lab New York

March 27, 2020

Jesica Maria Ramirez Toscano

0.0.1 Part 1: Variable Creation

```
[164]: import pandas as pd
import numpy as np
arrests = pd.read_csv('arrests.csv')
demo = pd.read_csv('demo.csv')
demo['bdate'] = pd.to_datetime(demo['bdate'], utc=False)
arrests['arrest_date'] = pd.to_datetime(arrests['arrest_date'], utc=False)
```

1. We filter the arrest to the ones that occurred post-implementation.
2. Since we need information about past arrests and potential felony re-arrests, we merge the post-arrests with the total arrests by person_id. So each arrest will be linked to a post-arrest of the same individual. > Note: **arrest_post** refers to the data of arrests post-implementation. **tr** refers to the merged data of arrests_post with all the arrests. So each arrest in this data set is linked to a post-arrest of the same individual.

```
[168]: arrests_post = arrests[arrests['arrest_date'] >= '2010-01-01'].copy()
tr = pd.merge(arrests,
              arrests_post.rename(columns={'arrest_date': 'date_post',
                                           'arrest_id': 'aid_post',
                                           'law_code': 'code_post'}),
              on='person_id')
```

3. We create different tables to obtain the number of prior misdemeanor arrests and felony arrests in the last 2 years and 6 months.

```
[169]: twoyear = tr[(tr['arrest_date'] >= tr['date_post'] - pd.DateOffset(years=2)) &
    ↳ (tr['arrest_id'] != tr['aid_post'])]
sixmonth = tr[(tr['arrest_date'] >= tr['date_post'] - pd.DateOffset(months=6)) &
    ↳ (tr['arrest_id'] != tr['aid_post'])]
twoyear = twoyear.groupby(['aid_post', 'law_code']).size().unstack().
    ↳ reset_index().fillna(0)
twoyear.rename(columns = {'aid_post': 'arrest_id', 'felony': 'fel_2y',
    ↳ 'misdemeanor': 'mis_2y'}, inplace=True)
```

```
sixmonth = sixmonth.groupby(['aid_post', 'law_code']).size().unstack().
    →reset_index().fillna(0)
sixmonth.rename(columns = {'aid_post': 'arrest_id', 'felony': 'fel_6m',
    →'misdemeanor': 'mis_6m'}, inplace=True)
```

So for the table **twoyear**, we have the post_arrests variable with the number of prior felony and misdemeanor arrests in the last two years.

```
[172]: twoyear
```

```
[172]: law_code arrest_id fel_2y mis_2y
0      000192be      0.0    1.0
1      000316a6      1.0    3.0
2      0005e711      1.0    3.0
3      000fca1a      1.0    3.0
4      00102ecf      0.0    3.0
...      ...      ...    ...
19068   fff23454      1.0    4.0
19069   fff2c58f      0.0    3.0
19070   fff4f37a      1.0    0.0
19071   fff50975      1.0    0.0
19072   fffdd268      1.0    1.0
```

```
[19073 rows x 3 columns]
```

4. To create the felony re-arrest binary variable, we need information about the potential future felony arrest of that individual. So first, we create a table called **year_ahead** using the **tr** dataset.

```
[173]: year_ahead = tr[(tr['arrest_date'] >= tr['date_post']) & (tr['arrest_id'] !=
    →tr['aid_post'])]
year_ahead = year_ahead[year_ahead['arrest_date'] <= year_ahead['date_post'] +
    →pd.DateOffset(years=1)]
year_ahead = year_ahead.groupby(['aid_post', 'law_code']).size().unstack().
    →reset_index().fillna(0)
year_ahead.rename(columns = {'aid_post': 'arrest_id', 'felony': 'felony_arrests'
    →}, inplace=True)
```

```
[174]: year_ahead[['arrest_id', 'felony_arrests']]
```

```
[174]: law_code arrest_id felony_arrests
0      000192be      0.0
1      000fca1a      0.0
2      00102ecf      0.0
3      001a4d40      0.0
4      002183f7      0.0
...      ...      ...
```

8975	ffd580fe	0.0
8976	ffe20b13	1.0
8977	fff23454	0.0
8978	fff2c58f	0.0
8979	fffd268	0.0

[8980 rows x 2 columns]

With this table, we can create a binary variable of re_arrest (1 if the individual has one or more felony arrests during one year following the arrest, 0 if the individual has no felony re-arrest)

```
[176]: year_ahead['re_arrest'] = np.where(year_ahead['felony_arrests'] > 0,1,0)
```

5. With twoyear, sixmonth, year_ahead tables, we can now fill the data in arrests_post about the number of prior felony arrests and misdemeanor arrests in the last 2 years and 6 months, and the binary variable re_arrest (felony re-arrest).

```
[177]: arrests_post = arrests_post.merge(twoyear, on='arrest_id', how='left').fillna(0)
arrests_post = arrests_post.merge(sixmonth, on='arrest_id', how='left').fillna(0)
arrests_post = arrests_post.merge(year_ahead[['arrest_id', 're_arrest']],
    on='arrest_id', how='left').fillna(0)
```

```
[178]: arrests_post
```

```
[178]:
```

	person_id	arrest_date	arrest_id	law_code	fel_2y	mis_2y	fel_6m	\
0	14	2010-01-01	e6a8cdb3	felony	1.0	0.0	1.0	
1	350	2010-01-01	be6e57e3	misdemeanor	1.0	5.0	0.0	
2	1409	2010-01-01	8af55340	misdemeanor	0.0	2.0	0.0	
3	2850	2010-01-01	cbb41c9a	misdemeanor	1.0	2.0	0.0	
4	2945	2010-01-01	65c81707	misdemeanor	2.0	2.0	1.0	
...	
21514	16253	2011-12-31	93ad5fb5	felony	1.0	1.0	1.0	
21515	16287	2011-12-31	fee2d3a2	misdemeanor	0.0	2.0	0.0	
21516	17822	2011-12-31	e94fc5a0	misdemeanor	0.0	0.0	0.0	
21517	18044	2011-12-31	86e5491e	felony	0.0	0.0	0.0	
21518	19164	2011-12-31	032164a7	misdemeanor	0.0	0.0	0.0	

	mis_6m	re_arrest
0	0.0	1.0
1	3.0	0.0
2	1.0	0.0
3	2.0	0.0
4	1.0	0.0
...
21514	0.0	0.0
21515	1.0	0.0

```

21516    0.0    0.0
21517    0.0    0.0
21518    0.0    0.0

```

[21519 rows x 9 columns]

6. Finally, we include data about the home precinct, age, and gender of the individual in each arrest. >For the age variable, we obtain the difference in the arrest date and the birthdate (the result is in days, we convert it to years.)
For the gender variable, we noticed it has four unique values: M, F, male, female. So we changed male and female values as M and F.

```

[179]: final = pd.merge(arrests_post, demo, on='person_id')
final['age'] = ((final['arrest_date'] - final['bdate']) / np.timedelta64(1,
→ 'Y')).round().astype(int)
final.drop(['bdate', 'arrest_id'], axis=1, inplace=True)

```

```

[180]: final.gender.unique()

```

```

[180]: array(['M', 'F', 'male', 'female'], dtype=object)

```

```

[181]: final.loc[final['gender'] == 'male', 'gender'] = 'M'
final.loc[final['gender'] == 'female', 'gender'] = 'F'

```

```

[182]: print(final.gender.unique())
final

```

['M' 'F']

```

[182]:
      person_id  arrest_date  law_code  fel_2y  mis_2y  fel_6m  mis_6m  \
0             14  2010-01-01    felony    1.0    0.0    1.0    0.0
1             14  2010-09-28    felony    1.0    0.0    0.0    0.0
2            350  2010-01-01  misdemeanor    1.0    5.0    0.0    3.0
3            350  2010-08-23  misdemeanor    0.0    5.0    0.0    0.0
4           1409  2010-01-01  misdemeanor    0.0    2.0    0.0    1.0
...          ...          ...        ...    ...    ...    ...    ...
21514       13443  2011-12-31  misdemeanor    0.0    0.0    0.0    0.0
21515       14245  2011-12-31  misdemeanor    0.0    0.0    0.0    0.0
21516       17822  2011-12-31  misdemeanor    0.0    0.0    0.0    0.0
21517       18044  2011-12-31    felony    0.0    0.0    0.0    0.0
21518       19164  2011-12-31  misdemeanor    0.0    0.0    0.0    0.0

      re_arrest  gender  home_precinct  age
0             1.0     M              58   21
1             0.0     M              58   22
2             0.0     M              38   42
3             0.0     M              38   42

```

4	0.0	M	61	20
...
21514	0.0	M	12	23
21515	0.0	F	1	21
21516	0.0	M	66	18
21517	0.0	M	66	40
21518	0.0	M	64	19

[21519 rows x 11 columns]

0.0.2 Part 2: Statistical Analysis » Program Evaluation

```
[183]: import matplotlib.pyplot as plt
import statsmodels.api as sm
import seaborn as sns
from statsmodels.discrete.discrete_model import Probit
```

1. First, we import data about the treatment and control precincts
2. Then, we are only interested in measuring the effect of the program for the first time an individual receives treatment, we filter the data to the first arrest of each individual in the post-implementation period.

```
[184]: treat = pd.read_csv('treatment_assignment.csv')
treat.rename(columns={'precinct' : 'home_precinct'}, inplace=True)

first = final.groupby('person_id').agg({'arrest_date':min}).reset_index()
first = first.merge(final, on=['person_id', 'arrest_date'])
```

If we look at the data in the treatment_assignment data set, there are 30 precincts (control and treatment precincts), whereas in the data set of first arrests post-implementation period, there are 77 different precincts.

```
[185]: print('Treatment-control precincts: {}'.format(len(treat.home_precinct.
    ↪unique()))
print('Arrests post-implementation precincts: {}'.format(len(first.home_precinct.
    ↪unique()))
```

Treatment-control precincts: 30

Arrests post-implementation precincts: 77

In this sense, we have two options:

- a) Assume that the precincts not included in the treatment_assignment data set are also CONTROL.
- b) Assume that the treatment_assignment is complete and those precincts were chosen to study because they are similar to each other.

I'm going to follow the option b), and drop the observations that don't fall in the control and treatment precincts.

```
[186]: data_eval = first.merge(treat, on=['home_precinct'], how='right')
data_eval.drop(['person_id', 'arrest_date'], axis=1, inplace=True)
data_eval
```

```
[186]:
```

	law_code	fel_2y	mis_2y	fel_6m	mis_6m	re_arrest	gender	\
0	misdemeanor	1.0	0.0	0.0	0.0	0.0	F	
1	misdemeanor	1.0	1.0	0.0	1.0	0.0	M	
2	felony	0.0	1.0	0.0	0.0	0.0	M	
3	misdemeanor	1.0	3.0	1.0	3.0	1.0	M	
4	misdemeanor	0.0	4.0	0.0	0.0	0.0	F	
...	
3660	felony	1.0	2.0	0.0	1.0	0.0	M	
3661	misdemeanor	1.0	1.0	1.0	1.0	1.0	M	
3662	felony	0.0	1.0	0.0	0.0	0.0	M	
3663	misdemeanor	0.0	2.0	0.0	2.0	0.0	M	
3664	misdemeanor	0.0	0.0	0.0	0.0	0.0	M	

	home_precinct	age	treatment_status
0	73	25	control
1	73	31	control
2	73	27	control
3	73	31	control
4	73	36	control
...
3660	74	31	treatment
3661	74	19	treatment
3662	74	24	treatment
3663	74	29	treatment
3664	74	24	treatment

[3665 rows x 10 columns]

- Before evaluating the success of the program, we change the values of the following variables:

- gender to 1 for Men and 0 for Female
- treatment_status to 1 for treatment and 0 for control
- law_code to 1 for felony and 0 for misdemeanor

```
[187]: data_eval['gender'] = np.where(data_eval['gender']=='M', 1, 0)
data_eval['treatment_status'] = np.where(data_eval['treatment_status']=='treatment', 1, 0)
data_eval['law_code'] = np.where(data_eval['law_code']=='felony', 1, 0)
```

- To analyze the effectiveness of this program, we regress the re_arrest variable on the rest of the covariates. > Since the dependent variable is binary, we must estimate heteroscedasticity robust standard errors

```
[193]: IND_VARS = ['treatment_status', 'age', 'gender', 'law_code', 'fel_2y', 'mis_2y', 'fel_6m', 'mis_6m']
all_ = sm.add_constant(data_eval[IND_VARS])
model1 = sm.OLS(data_eval['re_arrest'], all_).fit(cov_type='HC1')
model1.summary()
```

```
[193]: <class 'statsmodels.iolib.summary.Summary'>
"""
                                OLS Regression Results
=====
Dep. Variable:                re_arrest    R-squared:                0.586
Model:                        OLS          Adj. R-squared:           0.585
Method:                      Least Squares  F-statistic:             131.7
Date:                        Thu, 26 Mar 2020  Prob (F-statistic):      9.18e-195
Time:                        22:50:14       Log-Likelihood:          340.69
No. Observations:            3665          AIC:                   -663.4
Df Residuals:                3656          BIC:                   -607.5
Df Model:                    8
Covariance Type:             HC1
=====
=====
coef      std err          z      P>|z|      [0.025
0.975]
-----
----
const      0.0136      0.017      0.800      0.424      -0.020
0.047
treatment_status  0.0035      0.007      0.474      0.635      -0.011
0.018
age         0.0003      0.001      0.570      0.568      -0.001
0.002
gender     -0.0005      0.009     -0.061      0.952      -0.018
0.017
law_code   -0.0082      0.008     -1.066      0.286      -0.023
0.007
fel_2y      0.0129      0.008      1.625      0.104      -0.003
0.028
mis_2y     -0.0054      0.005     -1.131      0.258      -0.015
0.004
fel_6m      0.5299      0.019     27.738      0.000      0.492
0.567
mis_6m     -0.0056      0.007     -0.803      0.422      -0.019
0.008
=====
Omnibus:            630.272    Durbin-Watson:           2.001
Prob(Omnibus):      0.000    Jarque-Bera (JB):        3726.130
Skew:              -0.684    Prob(JB):                0.00
```

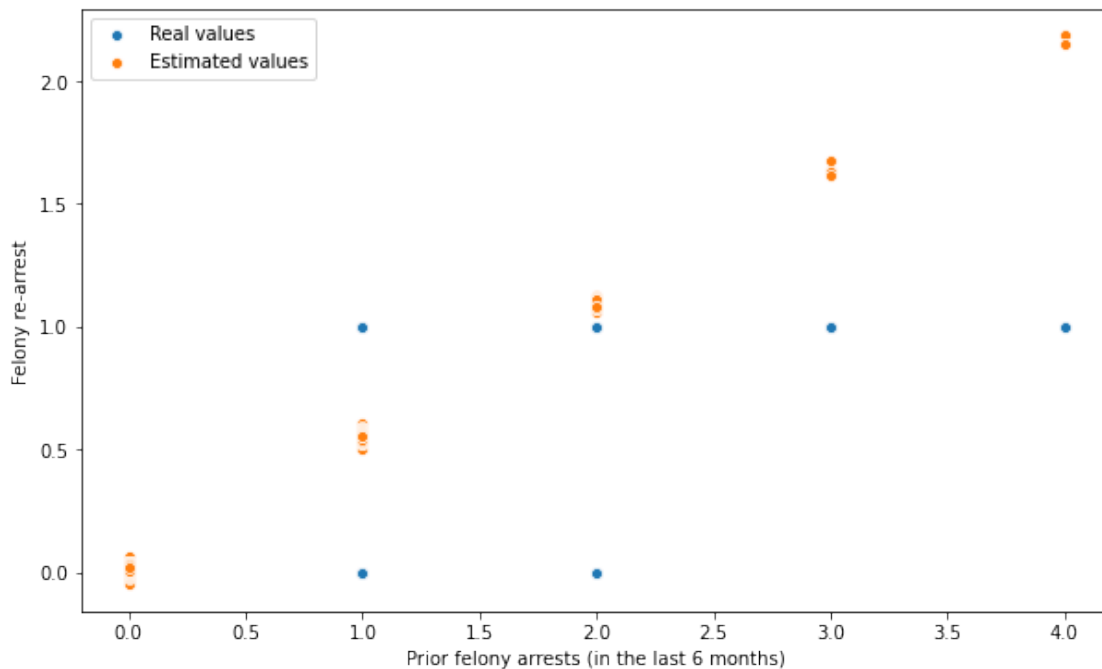
Kurtosis: 7.746 Cond. No. 145.
=====

Warnings:

```
[1] Standard Errors are heteroscedasticity robust (HC1)
"""
```

In this linear probability model, we observe that the treatment_status is not statistically significant, which may imply that there is no evidence to say that the program reduced or even affected the probability of felony re-arrest. In fact, the only variable that significantly explains the variation in the re_arrest probability is the recent-history (prior 6 months) of felony arrests. We plot this variable (prior felony arrests in the last 6 months) with the binary variable re_arrest and the estimated probability values of this model. In the graph above, we observe that the some estimated values are above one, and below zero (which makes no sense in probability).

```
[194]: plt.figure(figsize=(10,6))
sns.scatterplot(data_eval['fel_6m'],data_eval['re_arrest'], label='Real values')
sns.scatterplot(data_eval['fel_6m'],model1.fittedvalues, label='Estimated_
↪values')
plt.xlabel("Prior felony arrests (in the last 6 months)")
plt.ylabel("Felony re-arrest")
plt.show()
```



Looking to the graph above, we might agree that the independent variables and re-arrest appropriate model may not be linear. In this sense, we can use a probit model to estimate the effects of the independent variables on re-arrest probability.


```
[197]: probitm = Probit(data_eval['re_arrest'], all_).fit()
probitm.summary()
```

```
Optimization terminated successfully.
      Current function value: 0.148459
      Iterations 8
```

```
[197]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

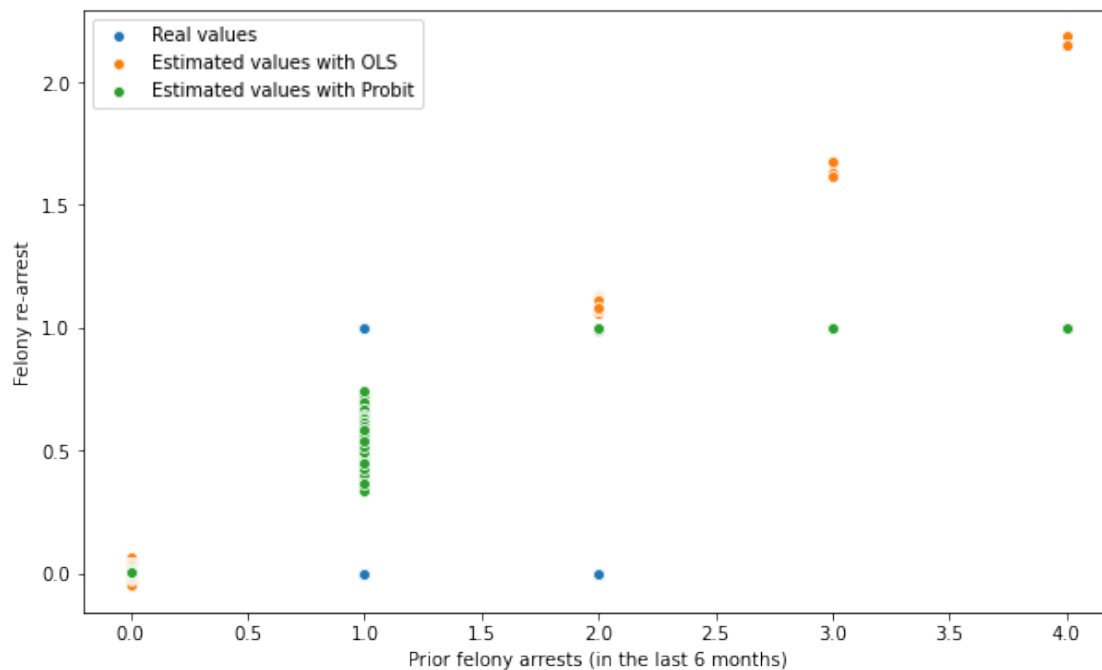
                                Probit Regression Results
=====
Dep. Variable:                  re_arrest      No. Observations:   3665
Model:                            Probit      Df Residuals:           3656
Method:                           MLE        Df Model:              8
Date:                Thu, 26 Mar 2020      Pseudo R-squ.:         0.6264
Time:                22:58:58              Log-Likelihood:        -544.10
converged:                      True        LL-Null:             -1456.5
Covariance Type:            nonrobust      LLR p-value:           0.000
=====
=====
      coef      std err          z      P>|z|      [0.025
0.975]
-----
----
const          -2.5018      0.206    -12.146      0.000     -2.906
-2.098
treatment_status -0.0204      0.089     -0.229      0.819     -0.195
0.154
age             0.0053      0.007      0.801      0.423     -0.008
0.018
gender         -0.0223      0.104     -0.213      0.831     -0.227
0.182
law_code       -0.0479      0.110     -0.434      0.664     -0.264
0.168
fel_2y          0.1139      0.086      1.325      0.185     -0.055
0.282
mis_2y         -0.0578      0.053     -1.083      0.279     -0.162
0.047
fel_6m          2.5577      0.122     20.924      0.000      2.318
2.797
mis_6m         -0.0764      0.069     -1.112      0.266     -0.211
0.058
=====
=====
      """
```

In this probit model, again, the treatment_status has no impact on the re-arrest variable. In this sense, there is no evidence that the program had an impact on the felony re-arrest probability. Most of the variation

of the re-arrest probability is explained by past felony arrests in the last 6 months, which in this model specification also shows significance.

Now, in the graph above, we look at prior felony arrests in the last 6 months with the binary variable *re_arrest* and the estimated probability values of the OLS and probit models. The estimated probability with the probit model is bounded within 1 and 0.

```
[199]: plt.figure(figsize=(10,6))
sns.scatterplot(data_eval['fel_6m'],data_eval['re_arrest'], label='Real values')
sns.scatterplot(data_eval['fel_6m'],model1.fittedvalues, label='Estimated values_
→with OLS')
sns.scatterplot(data_eval['fel_6m'],probitm.predict(all_), label='Estimated_
→values with Probit')
plt.xlabel("Prior felony arrests (in the last 6 months)")
plt.ylabel("Felony re-arrest")
plt.show()
```



The results show no evidence of the program being succesful in reducing felony re-arrest in a 12-month follow-up. Adittionally, these results may be robust given that we have set several control variables such as age, gender, past arrests and also, assuming that the RCT was done correctly (we could verify this with balance tests).